Mira: AI-Powered Project Intelligence Assistant

Sathish Doraiswamy | Nexora Pvt. Ltd.

Applied Agentic AI for PMs/TPMs, capstone project

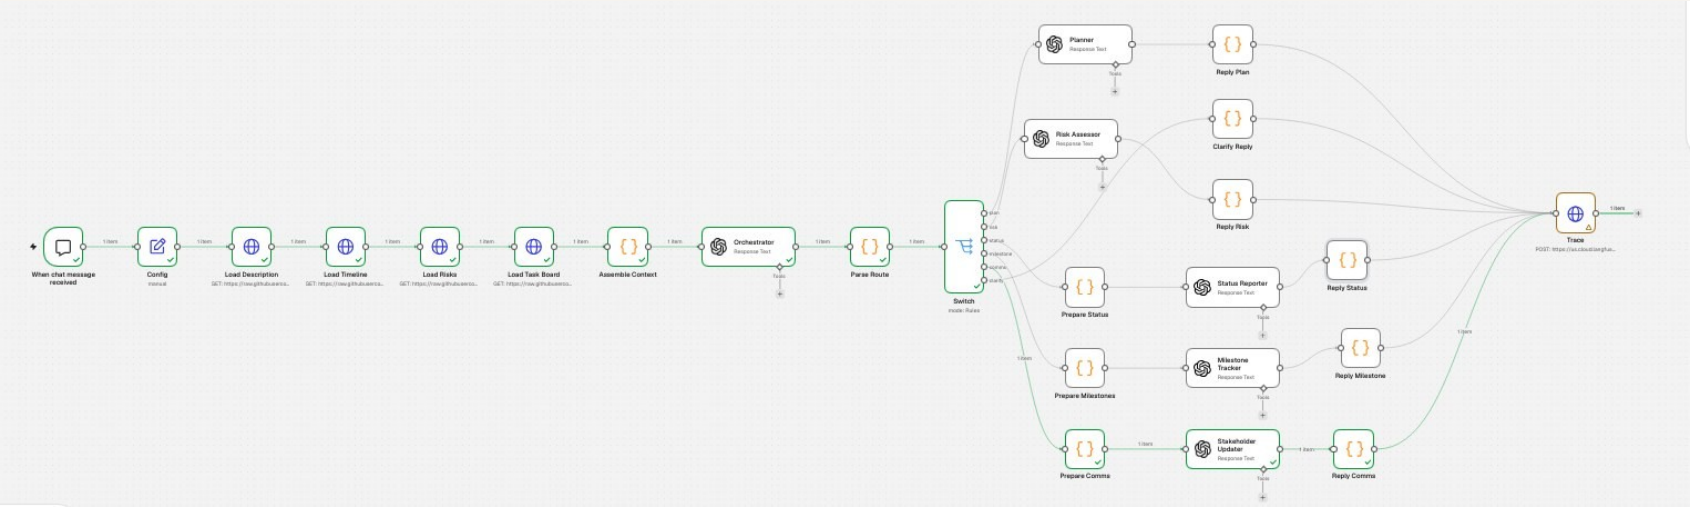
PMs lose hours to plans and reports, and the biggest AI risk is invention

- PMs manage several projects at once with little tool support
- A project plan takes 3 to 4 hours; a weekly status report takes 1 to 2 hours
- Risk assessments come in inconsistent formats
- Missed milestones are found in retrospectives, not in time
- Biggest AI risk: invented milestones, risks or task statuses

Code computes the facts, the model only writes the narrative

- An Orchestrator classifies each request and checks there is enough data
- Five specialists: Planner, Risk Assessor, Status Reporter, Milestone Tracker, Stakeholder Updater
- Code nodes compute counts, dates and filters, so numbers are never guessed
- Vague requests get a fixed reply asking for scope, goals, timeline and team, with no AI call

A router sends each request to one specialist



- Router / Dispatcher: the Switch sends each request to one of six branches
- Every branch ends in one Langfuse Trace node (rightmost)

Four tools, each chosen for a reason

Tool	Role	Why
n8n Cloud	Workflow	Visual router, CSV loading, Code nodes for exact counts
OpenAI gpt-4o-mini	LLM	Cheap and fast for structured writing from supplied data
Langfuse	Observability	One trace per request with input, output and route
GitHub raw files	Data	Four project files, no authentication needed

All 12 tests pass with zero invented items

12 of 12

tests pass

0

hallucinated items in 12 traced runs

100%

task counts accurate

Task board: 5 Done, 3 In Progress, 1 Blocked, 16 To Do. All 4 vague inputs asked for more detail instead of guessing.

Tests exposed six problems, and each got a fix

Test	What went wrong	Fix
T1	Plan missed the Week 1 start-date assumption	Prompt rule added
T3	Unrequested top 3 risks printed	Answer only what was asked
T7	Top 3 risks generic and inconsistent	Rank by task evidence
T10	Blocker and To Do list missing	Pass description through; list every status
T12	Overdue task missing under Risks	Rule added
Tracing	Legacy endpoint showed nothing	Moved to OTLP; added Final Reply node

A typical user costs about 20 cents a month

$0.00008

vague request (Orchestrator only)

~$0.001

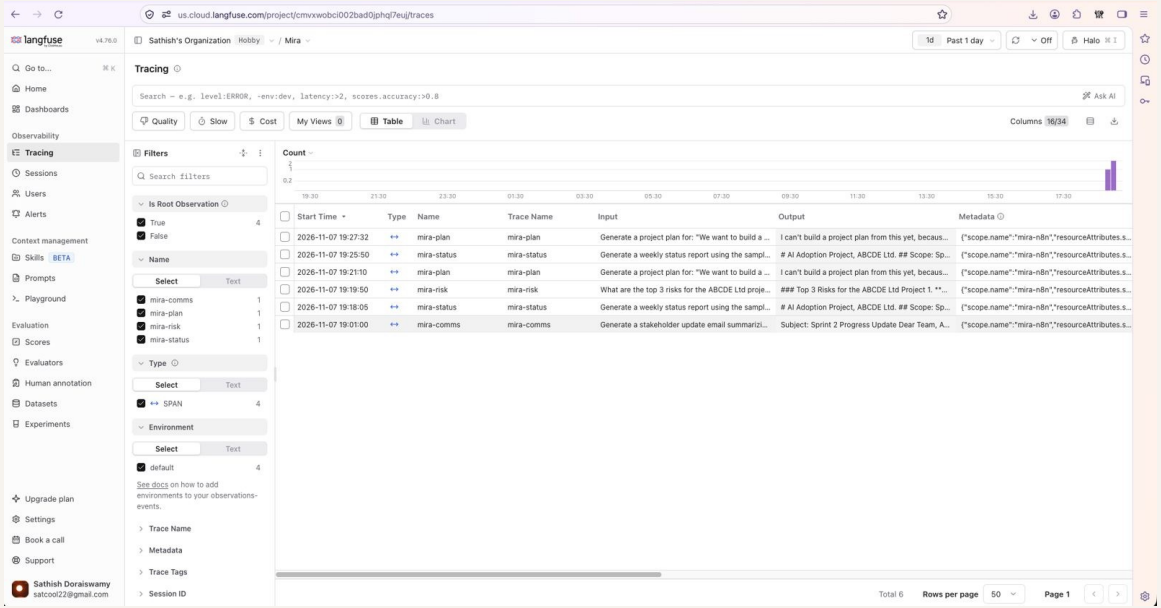
heaviest request: status report, 25 tasks

~$0.20

per user per month

A typical day of 15 requests costs about $0.0093 (gpt-4o-mini, $0.15 per 1M input and $0.60 per 1M output tokens).

Every request leaves one trace, tagged by route



- Input, output and route per request
- Routes: mira-plan, risk, status, milestone, comms
- Used to find and fix the failures

Rollout: pilot with 3 PMs, then widen behind a quality gate

- Weeks 1 to 4: pilot with 3 PMs, every output reviewed
- Weeks 5 to 8: expand to 10 PMs, live task board, approval step for emails
- Week 9 onward: monthly metric review
- Gate to widen: 0 fabricated items reaching a stakeholder, status accuracy 100%, satisfaction 4 of 5 or higher

Lesson: compute facts in code and let the model explain

- The loop that worked: evaluate, find the failure, fix the prompt, re-run
- Next: lower temperature for steadier answers
- Next: code-ranked risks, automated regression tests
- Next: live Trello connection and a human approval step